



CallGuard AI

FOR FCA-REGULATED FIRMS, APPOINTED REPRESENTATIVES & PRINCIPALS

Audit-ready evidence from every customer conversation.

CallGuard AI transcribes and scores 100% of your firm's calls against your own compliance scorecard — every verdict backed by the exact quote from the transcript. When the FCA, your auditor or the Ombudsman asks you to show your work, the answer is an export, not a six-month file review.

100%

of calls transcribed, scored and evidenced — not a 5% sample

Minutes

from call ending to a scored result, versus weeks for manual review

Zero

customer personal data reaching the AI — redacted before it is stored

THE PROBLEM

You are accountable for every conversation. You can evidence about five per cent of them.

Consumer Duty asks firms to demonstrate good outcomes systematically, not anecdotally. The strengthened appointed representative regime asks principals to evidence real oversight of firms they do not own. Both expect proof that manual file sampling cannot produce — and the gap between what you are accountable for and what you can show is where regulatory risk now sits.

5–10%

CALLS REVIEWED TODAY

The typical sample rate for manual file checking. The other 90% are never heard.

60–90 min

PER FILE REVIEWED

Senior reviewer time per advice file — your scarcest and most expensive resource.

Weeks

UNTIL A BREACH SURFACES

By which point the recommendation has gone out and the customer has acted on it.

MANUAL SAMPLING TODAY

- ✗ A single-digit sample; the rest are invisible
- ✗ Scoring varies between reviewers, and between weeks
- ✗ Findings land days or weeks after the advice was given
- ✗ Vulnerability signals slip past unflagged
- ✗ Outcomes evidenced by anecdote and spreadsheet

WITH CALLGUARD AI

- ✓ Every call transcribed, scored and filed as evidence
- ✓ One consistent standard, calibrated to your interpretation
- ✓ Critical breaches flagged within minutes — or mid-call
- ✓ Vulnerability indicators screened on every conversation
- ✓ Outcomes evidenced systematically, with quotes

THE QUESTION BEHIND ALL OF IT

Every firm can describe its compliance process. Far fewer can produce, on demand, the evidence that it happened on a specific call, with a specific customer, on a specific date. **That evidence is what CallGuard AI produces — automatically, for every conversation.**

HOW IT WORKS

Calls in. Compliance evidence out.

CallGuard sits alongside the telephony you already use. There is no rip-and-replace, no new dialler, and no integration project — whether your team uses a cloud dialler, a compliance recorder, Microsoft Teams, or simply uploads recorded advice files.

1 Bring in your calls, however they arrive

Drag-and-drop upload, automated collection from your recording server, a webhook from your dialler, or our REST API. CloudTalk and Zoho CRM connect natively; Twilio and AWS Connect stream live. Audio is encrypted the moment it lands.

2 Score against **your** scorecard, not a generic one

Upload your existing QA scorecard — the one your compliance team already uses. Add your suitability rules, product guides and call standards to the knowledge base. Every call is transcribed with the adviser and customer separated, then judged criterion by criterion against your firm's standard.

3 Act on it, and keep the proof

Pass or fail per criterion with a supporting quote, a weighted overall score, a breach register, and a coaching brief for the adviser. Critical breaches alert by email, Slack or webhook. Results write back to your CRM. Everything is retained as the audit record.

It learns your firm's interpretation, not a textbook's

Where the AI is uncertain, or where a consent step is borderline, the call is routed to a review queue rather than being scored automatically — no false pass, no false breach. When your compliance officer corrects a verdict, that correction becomes a worked example that shapes future scoring. Mark your best calls as exemplars and they set the reference standard. After a few dozen corrections, CallGuard scores the way your senior reviewer scores. **Your interpretation of "suitable", "fair value" and "good outcome" is the one that counts.**

THE PLATFORM

Five things CallGuard does that a QA spreadsheet cannot.



Every call scored, with the proof attached

Not a sample — every conversation. Each criterion on your scorecard returns a pass or fail, a confidence level, and the exact sentence from the transcript that justifies it. You are never asked to take the AI's word for it; you can read the line it relied on.



We score the sale, not just the call

A single recommendation usually spans three or four conversations. CallGuard joins them into one customer journey and assesses compliance across the whole thing — so a disclosure made on the second call still counts, and one missed everywhere is caught. That is how the regulator looks at an outcome, so it is how we score it.



Breaches caught mid-call, not next week

For teams that stream live, CallGuard listens as the conversation happens and flags a critical compliance failure before the adviser hangs up — only when it is confident, so supervisors are not swamped with noise. A breach prevented costs nothing to remediate.



Customer personal data never reaches the AI

Names, addresses, card numbers and health details are stripped out at the moment of transcription and replaced with typed tags. They are never stored in our system and never sent to a language model. It is the first question your DPO will ask, and it is designed into the pipeline rather than bolted on afterwards.



The audit pack, on demand

When the FCA, your auditor, your principal or the Ombudsman asks you to show your work, export it: every call, every verdict, every supporting quote, every correction your compliance team made, and who made it. What used to be a file-review project becomes a download.



FOR PRINCIPAL FIRMS

You carry the regulatory risk for conversations happening in firms you don't own.

A principal is responsible for the regulated activity of every appointed representative it accepts. Since the regime was strengthened, that comes with an annual review of each AR and an annual self-assessment of your oversight arrangements. Most principals meet it with periodic file requests, an annual visit and a great deal of trust. CallGuard replaces the trust with evidence.

Consistent oversight across the network

One scorecard applied identically to every AR, adviser and call — so comparing two ARs means something. No reviewer variance, no differing local standards.

EVIDENCE FOR THE ANNUAL AR REVIEW

Risk ranked, so attention goes where it belongs

Adviser- and firm-level risk scoring surfaces the AR whose disclosure rates are slipping, before it becomes a complaint file. Oversight follows the data, not the rota.

TARGETS YOUR MONITORING PLAN

Separate portals, one standard

Each AR sees their own firm's calls, scores and coaching. You see everything, across the whole network, in one view. Data stays properly partitioned between representatives.

MULTI-FIRM BY DESIGN

A defensible record, not a recollection

Every verdict is quoted, dated, attributed and retained. When you write your self-assessment, the supporting evidence is already assembled rather than reconstructed from memory.

EVIDENCE FOR YOUR SELF-ASSESSMENT

FOR APPOINTED REPRESENTATIVES

The same platform works in the other direction. If your principal has asked how you monitor advice quality, CallGuard is the answer you can show them: 100% of your calls scored against their standards and available for inspection. **It turns your next oversight visit from a document hunt into a login.**

A worked example

A principal with 18 ARs and 60 advisers samples roughly 5% of calls. At 90 minutes a file, full coverage would need several more reviewers. With CallGuard all 60 are scored on every call, breaches surface the day they happen, and each AR's review evidence compiles itself. **Oversight stops scaling with the network.**

THE EVIDENCE

Every verdict shows its working.

This is what a scored criterion looks like in CallGuard. Two things to notice: the evidence is a direct quote, and the customer's personal details have already been replaced with typed tags before anything was stored or analysed.

Discloses that the policy has no cash-in value

FAIL · CRITICAL

EVIDENCE

"So that's [MONEY_1] a month and it covers you right through to [AGE_1] — shall I get that set up for you today, [NAME_GIVEN_1]?"

REASONING

The adviser moved from premium confirmation directly to closing. No statement was made at any point in the call that the policy has no surrender or cash-in value.

CONFIDENCE

0.93 — scored automatically. Below 0.70 this item would route to your review queue instead.

WHERE

Call 2 of 3 · 11:42 · Adviser: J. Whitfield · Customer journey #4821

Explores indicators of customer vulnerability

PASS

EVIDENCE

"You mentioned you'd been unwell recently — is now a good time to be going through this, or would you rather I called back when someone can sit with you?"

REASONING

The adviser identified a health disclosure, checked capacity to proceed, and offered a deferral with support present. Recorded against the firm's vulnerable customer standard.

Uncertainty goes to a human, not a guess

Where the AI is not confident — particularly on consent and affordability steps — the item is placed in a review queue for your compliance officer instead of being auto-scored. Their decision is recorded, and it teaches the model your standard.

Corrections are the audit trail too

Every override is logged with who made it, when, and why. That record is as useful in an oversight conversation as the original verdict: it shows a compliance function actively supervising its own controls.

SECURITY & DATA PROTECTION

Built for a firm that has to answer for its processors.

You are accountable for your data processors as well as your advisers. These are the answers your DPO, your IT reviewer and your principal will ask for — set out plainly so you can pass this page straight to them.

Personal data is removed before storage

Names, addresses, postcodes, phone numbers, card and bank details and health information are detected and replaced with typed tags at the point of transcription. The raw values are never written to our database and never form part of any AI prompt.

Audio encrypted at rest

AES-256-GCM encryption for every stored recording, with keys held separately from the data.

Encrypted in transit

TLS 1.2 or above for all connections, including dialler webhooks, live audio streams and CRM write-back.

UK data residency

Data is hosted in the UK, and speech processing is pinned to the EU region with model-improvement use explicitly opted out.

Role-based access, with an audit log

Administrator, supervisor, viewer and adviser roles; advisers see only their own calls. Two-factor authentication is available, and privileged actions are written to an immutable audit log.

Your retention policy, enforced automatically

Set the retention period your firm requires. Recordings and transcripts are purged on schedule without anyone having to remember, and deletion is recorded.

UK GDPR Article 28 DPA

A data processing agreement is available on request, and our sub-processor list is published publicly — so you know exactly who touches your data before you sign anything.

Your data is not used to train models

Improvement using anonymised customer-derived data is opt-in, off by default, and recorded with a timestamp if you ever choose to enable it.

PRICING

Published prices. No enterprise quote required.

A seat is any user account on the platform: the advisers whose calls we score, and the compliance officers and supervisors who review the results. Seats are billed per account, not per call. Volume discounts apply as the team grows, and there is no multi-year minimum.

Starter

From £199 / seat / mo

For firms replacing manual sample QA with scoring on every call.

- ✓ 100% of calls transcribed and scored
- ✓ Your own scorecard
- ✓ Evidence quotes per criterion
- ✓ Breach register and dashboards
- ✓ Email alerts on critical breaches
- ✓ Upload and REST API ingest

Growth **MOST FIRMS**

From £299 / seat / mo

For compliance teams who want the AI calibrated to their own interpretation.

- ✓ Everything in Starter
- ✓ Learning from your corrections
- ✓ Exemplar call library
- ✓ Per-adviser coaching memory
- ✓ Slack alerts and signed webhooks
- ✓ CRM write-back, including Zoho

Pro

From £399 / seat / mo

For principals, AR networks and firms with audit-grade requirements.

- ✓ Everything in Growth
- ✓ Journey (sale-level) scoring
- ✓ Multi-firm scorecards and portals
- ✓ Live mid-call breach detection
- ✓ Priority support
- ✓ DPA and security review support

No change to your phone system, and no separate transcription bill. If you can produce a recording, CallGuard can score it — by upload, from your recording server, or through your dialler's webhook. Live streaming is optional, not a prerequisite. Transcription, redaction, scoring and coaching are all included in the seat price, so there is no per-minute bill to forecast.

How does this work across an AR network?

Each appointed representative is set up as its own firm with its own users and data, and the principal gets a view across all of them. Seats are counted across the network, so the volume discount applies to the whole group rather than firm by firm.

Is there a minimum term?

No multi-year minimum. Annual billing attracts a discount, but you can start monthly and move to annual once the product has proved itself on your own calls.



GETTING STARTED

From first call to full coverage, in days.

No implementation project, no solution architect, no 90-day pilot — most firms are scoring live calls in the same week they decide to go ahead.

STEP ONE

A 15-minute call

Bring five of your own recordings. We score them live, against your scorecard, on your data.

STEP TWO

Set-up

We load your scorecard and compliance material, and connect your recordings or dialler.

STEP THREE

Calibration

Your compliance officer reviews the first scored calls and corrects anything that reads wrong.

THEN

Full coverage

Every call from that point forward is scored, evidenced and retained automatically.

What we need from you

Your existing QA scorecard, in whatever form it exists today. Any compliance manuals or call standards you want the AI to know. A way to reach your recordings. That is the list.

What we do not need

A change of dialler. A new recording platform. Your customers' personal data — it is redacted before it reaches us. Or a year of your compliance team's time to configure it.

BEFORE YOU CHOOSE ANY VENDOR

01

"Show me the evidence behind one verdict."

A score with no supporting quote is an opinion. If they cannot show the sentence that justified the verdict, it is not evidence you can put in front of a regulator.

02

"Where does our customers' personal data go?"

Ask whether names, card details and health information reach the AI model, and whether they are stored. The answer should be no on both counts, by design.

03

"How do you score a sale that spans four calls?"

Most tools score calls in isolation. Consumer Duty is about the customer's outcome, which rarely happens in one conversation.

Bring five of your calls. We'll score them live.

The fastest way to judge CallGuard is to watch it score your own conversations against your own scorecard — not a demo tenant, and not someone else's data.

EMAIL

hello@callguardai.co.uk

WEB

callguardai.co.uk