CallGuard AI

PERSONAL BRIEF FOR [NAME]

Every advice call. Scored against your firm's bar. By AI that learns from your compliance team.

[Name],

[Friend] mentioned you've been wrestling with the same compliance bottleneck most advice firms have right now: file checks at a single-digit sample, Consumer Duty raising the evidence bar, and a checking team that can't realistically scale to cover what the FCA actually expects.

This is the project [Friend's first name] and I have been building. CallGuard was designed inside the UK financial planning sector from day one — the original scorecards covered ATR, capacity for loss, vulnerable client screening, COBS 9.2 suitability, charges disclosure and Consumer Duty outcomes. We've moved a long way in the last few weeks. I wanted to send this through directly rather than via a generic deck.

Five pages. Worth a 15-minute call if anything resonates.

— Kofi

THE PROBLEM

You're checking 5% of advice files. Consumer Duty expects more.

Most advice firms run a manual file-check at a 5-10% sample. Each file takes a senior compliance reviewer 60-90 minutes. Variation between reviewers is meaningful. And the FCA, post-Consumer Duty, has made it clear that "fair value" and "good outcomes" need systematic evidence — not a sampled spot-check three months after the recommendation went out.

5–10%

FILES REVIEWED TODAY

Industry average sample rate for manual file checking inside an advice business

90 min

PER FILE REVIEW

Senior reviewer time per advice file (recordings, fact-find, suitability letter, evidence)

s166

THE RISK YOU CAN'T SEE

Skilled-person reviews land on firms whose monitoring is anecdotal, not systematic

STATUS QUO

- ✕ Single-digit sample, recommendations are mostly invisible
- ✕ Subjective scoring varies between checkers
- ✕ Reviews land weeks after the recommendation went out
- ✕ Vulnerable client signals slip past unflagged
- ✕ Consumer Duty outcomes are anecdotal, not evidenced
- ✕ DB transfer specialist sign-off can't be re-checked at scale

WITH CALLGUARD

- ✓ 100% of advice calls transcribed and scored automatically
- ✓ One AI scorer calibrated to your firm's interpretation
- ✓ Critical breaches surface within seconds of the call
- ✓ Vulnerability indicators flagged on every call
- ✓ Consumer Duty outcomes evidenced systematically
- ✓ Pension transfer scoring with specialist-pathway gates

HOW IT WORKS

Calls in. Suitability evidence out.

CallGuard sits alongside your existing recording infrastructure. Whether your advisers record on a desk dialer, on Teams, or via a dedicated compliance recorder, audio flows in, gets scored against your firm's scorecard, and surfaces in a dashboard your compliance director actually wants to open.

1

Ingest — live or batch

Drag-and-drop, REST API, SFTP/S3 sync for archived recordings, or live WebSocket streaming from a Twilio, AWS Connect or compliance recorder. Encrypted at rest with AES-256-GCM.

2

Score against your scorecard

Deepgram transcribes with speaker separation. Claude scores against your firm's scorecard, factoring in your KB (firm policies, product cheatsheets, FCA references), prior compliance corrections, and exemplar calls.

3

Surface what matters

Pass/fail per criterion, weighted overall score, exact quote evidence, an adviser-specific coaching brief, and a breach register. Critical fails alert your compliance lead instantly.

Scoring built around the FCA framework

Suitability (COBS 9.2)

ATR explored properly (not just a closed question), capacity for loss assessed, time horizon confirmed, existing arrangements reviewed, recommendation linked back to stated objectives.

Vulnerable customers (FG21/1)

Indicators screened for: bereavement, cognitive impairment, financial distress, time pressure. Vulnerability flag triggers the firm's protocol; absence of screening is itself a flag.

Charges & Consumer Duty

Combined monetary illustration (not just %), product charges, ongoing service fees, fair value test, plain-English explanation, customer asked to confirm understanding.

Critical fails (regulator-relevant)

Urgency / pressure language ("offer closes today"), unauthorised pension transfer steers without specialist sign-off, missing recording disclosure, undisclosed conflicts of interest.

RECENT BUILD

Score the call while the adviser is still on it.

The shift in the last month is platform-level. CallGuard isn't a "drag-and-drop your recordings" tool any more. The AI reads the call as it's happening, calibrates itself to your compliance team's interpretation, and writes the brief your compliance director would otherwise spend a Sunday afternoon producing.

📡 Live streaming + mid-call breach detection

Audio streams live from your dialer or compliance recorder. Every 30 seconds the AI scans the rolling transcript and flags critical breaches via signed webhook to your supervisor. The compliance lead can step in *during* the call, not three weeks after.

🎯 Confidence-gated alerts

The live scorer only emits a breach when the AI is >75% confident the criterion has clearly already failed. Designed to avoid alert fatigue and the "boy who cried wolf" problem from day one.

🧠 Per-firm AI learning

Your compliance officer corrects the AI on borderline calls. Each correction becomes a few-shot example feeding the next batch of scoring. After ~50-100 corrections the AI is calibrated to your firm's interpretation of "suitable" and "fair value," not a generic baseline.

⭐ Firm exemplar library

Mark gold-standard advice calls. The AI uses them as the reference bar when scoring future calls — your firm's "good" defines the AI's "good," not someone else's.

📊 AI Insights — written for your CO

A weekly strategic brief: what changed this week, the patterns worth attention, and 3-4 specific actions ("escalate Adviser X this week, run a team-wide refresher on capacity for loss, update the charges script to include £ illustration"). Written by Claude, not boilerplate.

🤝 Adviser coaching memory

Coaching builds on what was said to that adviser last week. If they've improved on the flagged area, the AI acknowledges it. If they haven't, it escalates the language. Personalised, not template.

What's coming next: live in-call coaching

Beyond breach detection: a real-time AI coach that whispers context-aware nudges to the adviser mid-call ("ask about capacity for loss now," "soften the close") via earpiece or on-screen prompt. Personalised to each adviser using the prior coaching memory above. No legacy QA vendor delivers this on modern LLMs.

INSIDE AN ADVICE FIRM

Five places it lands immediately.

COMPLIANCE TEAM

File checking at 100%, triaged for human review

AI scores every call; team focuses senior-reviewer hours on the calls flagged amber/red, not a random 5% sample. Reviewer load drops; coverage approaches 100%.

PENSION TRANSFER OVERSIGHT

DB transfer specialist sign-off, evidenced systematically

Every transfer-relevant call scored against the DB pathway: triage warning indicators, capacity for loss, sophisticated investor caveats, specialist signoff trail.

VULNERABLE CUSTOMER PROTOCOL

FG21/1 indicators flagged on every call, not by chance

Bereavement, cognitive, financial distress and time pressure markers detected systematically. Triggers the firm's vulnerability protocol within minutes, not weeks.

CONSUMER DUTY EVIDENCE

Outcomes data the FCA actually wants

Aggregate scoring data answers "are advisers giving fair-value advice consistently?" with hard numbers, not a sample-of-30 written narrative buried in last quarter's MI pack.

ADVISER DEVELOPMENT

Per-adviser coaching that compounds

Each adviser's coaching brief is built on top of last week's. The struggling 5% get tracked weekly; the strong performers get exemplar-tagged and feed the AI's reference bar.

SENIOR MANAGERS (SM&CR)

Reasonable steps you can actually demonstrate

Systematic monitoring evidence on file. If a regulator asks how you supervise outcomes, you point at the dashboard, not the spreadsheet. Reduces s166 / past business review exposure.

What I'd want from you, [Name].

Honest reaction first. If it lands, I'd love to run a small design-partner pilot inside [Firm] — bring 50-100 of your real (anonymised) advice calls, get them scored against your existing scorecard, and let your compliance team correct anything the AI gets wrong. Two weeks, no commitment. The corrections you make calibrate the AI to your firm's bar, so even the pilot data has standalone value to you.

DIRECT

kofi@properleads.co.uk

15-MIN WALKTHROUGH

Reply to the message [Friend] sent — pick a slot.

UK-based · Hosted in EU/UK regions · GDPR-compliant by default · DPA available · AES-256-GCM encryption at rest · JWT auth · role-based access